

x86 Real/Protected Mode Programming Tips for MEMOS

CS552 – Operating Systems
02/27/2020

Sasan Golchin

Today's agenda

- MEMOS-1
 - x86 real-mode programming
 - Address translation @ link-/run-time
 - Useful BIOS services
 - Example
 - Linking a flat binary program
- MEMOS-2
 - x86 protected-mode programming
 - Jumping into C from assembly
 - GRUB and the Multiboot standard
 - Example
 - Linking an ELF program
 - GRUB configuration to run our kernel!

MEMOS-1: Overview

- It's a Master Boot Record (MBR) program
 - Located at the first sector of the bootable device (Disk, USB, ...)
 - Limited to 512 bytes in size ending in signature bytes 0x55 and 0xAA
 - The signature tells the firmware that it's a valid runnable MBR, not some random data!
 - Copied into a conventional address (0x7C00) by the firmware (BIOS, UEFI)
 - When finished initializing the system, BIOS makes a jump to 0x7C00
 - Your work starts here!
- Your MBR's objective
 - Detect the memory layout
 - Print it on the screen

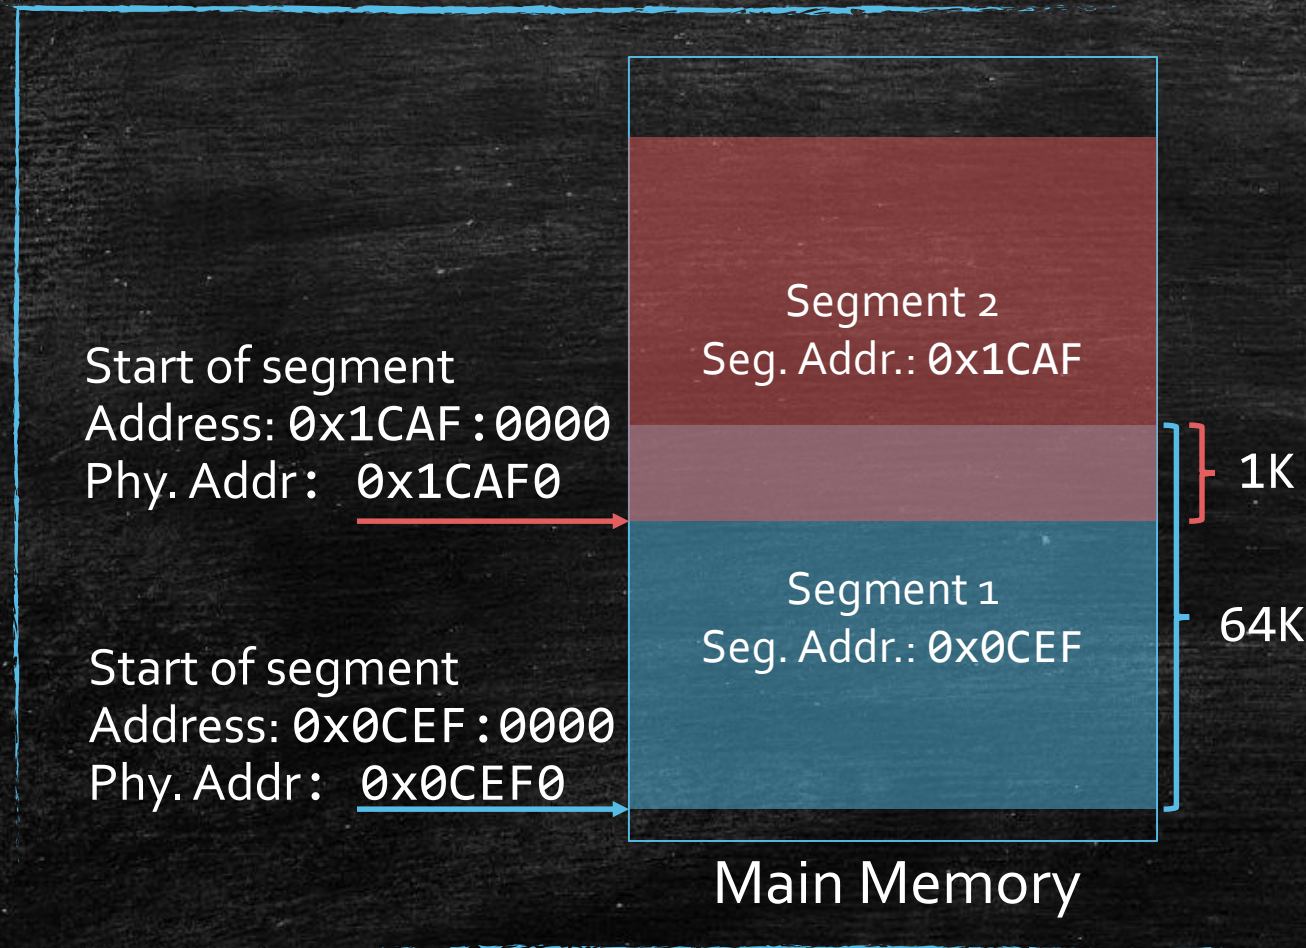
MEMOS-1: Where to begin?

- Need to know the state of the machine when MBR execution begins
 - Kind of instructions we can execute: 16-bit Real-Mode Code
 - How to address the memory: Real-Mode Segmentation
- Execution environment
 - No underlying Operating System – We're on our own! Wait! It gets even worse!
 - No library functions e.g. printf, scanf, itoa and etc.
 - Only BIOS interrupt service routines
- MBR program organization
 - Flat binary: No specific format, a mixed bag of code and data
 - The firmware just makes a jump to 0x7C00 and expects runnable code there!
 - Usually written in Assembly! Why not C? ([x86 Assembly Guide](#))

MEMOS-1: x86 Real-Mode

- 16-bit Instructions and Registers
 - 16-bit registers and operands by default
 - AX, BX, CX, DX, SI, DI, BP, SP
 - 32-bit registers are still available: EAX, EBX,...
- 20-bit Memory Address Space (Up to 1MB)
 - Direct access to physical memory
 - Firmware, user and devices co-exist – No memory protection!
 - Need for segmentation to access beyond 64KB
 - CS, DS, SS, ES, FS, GS
 - Don't forget to set your segment registers!
 - Memory segments can overlap.

MEMOS-1: Real-Mode Memory Segmentation



Issue:

20 Address lines ($2^{20} = 1\text{MB}$)

16-bit registers ($2^{16} = 64\text{KB}$)

`mov $val, (%cx)`

Solution:

16-bit segment registers

`PAddr = Operand + SEG << 4`

`mov $val, (%fs:%cx)`

Special purpose seg. regs.:

CS, DS, SS

More info: OSDEV

MEMOS-1: What is a flat binary

- The Executable and Linkable formats (e.g. ELF):
 - Organizes Code and Data into different sections
 - .text: machine instructions
 - .data: global tables and variables
 - .rodata: constant data (e.g. string literals, const variables,...)
 - .bss: uninitialized data
 - Each section can be loaded at a different location in the memory
 - Some data sections (like .bss) don't exist in the file but the loader inflates them in the main memory
 - Can contain information about symbols and for debugging purposes
 - The entry point (which function is the main function) for executable binaries
- Flat binary doesn't have any meta-data: Mixed bag of code and data
 - Order: By the order of your code and the as specified in the linker script

MEMOS-1: Address translation

- What are symbols in a program?
 - From programmer's perspective: Names of variables, functions and labels
 - From machine's perspective: Absolute or relative addresses!
 - Relative like: `jmp 1b # IP -= num. of bytes to get to 1b:`
 - Absolute like: `leaw myvariable, %ax # %ax = a number but what?`
- Compiler does not replace variables with addresses
- Linker does! How does it know where our program will end up in the main memory at run-time?
- Answer:
 - The `.org` directive in assembly
 - Specified in the linker script

MEMOS-1: BIOS routines come to rescue

- The firmware sets up the system devices in order to execute MBR
- It also provides a basic set of device drivers and service routines
- You can access those services by issuing software interrupts
 - Pass the parameters through registers
 - Invoke a software interrupt using the INT instruction
- Find the list of standard BIOS interrupt service routines [here](#)
- Useful BIOS interrupts:
 - INT 0x10: Video services (a very basic graphics driver)
 - INT 0x15: System services (e.g. to get system parameters)

MEMOS-1: Sample MBR program (mbr.S)

```
.code16
.globl _start
_start:
#Initialize the data segment
movw $0x7C0, %dx
movw %dx, %ds
#Get the address of string
leaw msg, %si
movw msg_len, %cx
#Print the greeting string
1:
lodsb #Loads DS:SI into AL
movb $0x0E, %ah
int $0x10
loop 1b
#Print "0x"

movb $'0', %al
movb $0x0E, %ah
int $0x10
movb $'x', %al
movb $0x0E, %ah
int $0x10
hlt
msg:
.string "MemOS: System Memory is:"
msg_len: .word . - msg
#Put the MBR Signature
.org 0x1FE
.byte 0x55
.byte 0xAA
```


MEMOS-1: Sample MBR program

Now we must compile and link our program:

- We want a flat binary file as BIOS does not support ELF or EXE
- The first instruction is located at `0x7C00` in the main memory

But how?

MEMOS-1: Introduction to linker scripts

- The linker collects compiled object files, resolves references to symbols to addresses and builds a binary executable/linkable file
- The linker script helps the linker to understand:
 - The desired binary format (OUTPUT_FORMAT)
 - Examples are: flat binary, 32-bit ELF, 64-bit ELF
 - The target system architecture which is later used by the loader software to verify if the binary is compatible with the architecture (OUTPUT_ARCH)
 - The entry point to the executable program (ENTRY): i.e. The first function that should be called by the loader to start the program
 - Sections of the output file (SECTIONS):
 - Where to put the sections in the main memory (base address) <- Addr. Translation
 - Where to get the sections from (which section of which object file)

MEMOS-1: Sample MBR program (link.ld)

```
OUTPUT_FORMAT("binary")
```

```
ENTRY(_start)
```

```
SECTIONS {
```

```
    . = 0x0;
```

```
    .section1 : { *.o (.*); }
```

```
}
```


MEMOS-1: Sample MBR program (Makefile)

```
as --32 memos-1.S -o memos-1.o
```

```
ld -m elf_i386 -T memos-1.ld memos-1.o -o memos-1
```

```
dd bs=1 if=memos-1 of=disk.img count=512
```

```
qemu-system-i386 -m 32 -hda disk.img
```


MEMOS-2: Overview

- We use GRUB's MBR instead of our MEMOS-1
- BTW, GRUB is a bootloader that:
 - Enumerates system resources (such as amount of memory available)
 - Switches the CPU to 32-bit protected mode
 - Finds your kernel executable file (ELF) from the boot media
 - This time, our kernel can be big file and is not limited to 512 bytes 😊
 - Loads it at 0x100000 (1MB) in the main memory
 - Passes the system information to your kernel
 - According to some standard format called the Multiboot Standard
 - Jumps to the memory address 0x100000
- This time, we have to comply with GRUB, instead of BIOS!

MEMOS-2: Where to begin?

- Need to know the state of the machine when GRUB calls our code
 - 32-bit protected mode with segmentation
 - 2 flat segments of 4GB: Can run code and access data anywhere in 0 to 4GB
- Execution environment
 - We are in Ring-0: Most privileged level, hence, can do anything
 - No access to BIOS services – This time, we are totally alone ☹
- Program organization
 - GRUB expects an ELF binary with a multiboot header

MEMOS-2: x86 Protected Mode

- 32-bit instructions and registers

- Still can access smaller parts of a register:



- Can address up to 4GB of memory

- Through segmentation provided by [GDT](#) and LDT
 - Here, segment registers (CS, DS, SS and etc) point to entries in GDT/LDT
- Through virtual memory (Paging)
 - Not needed now
- Provides memory protection by restricting types of instruction you can run in a segment of certain privilege
 - 4 Different levels: Ring-0 (most privileged) to Ring-3 (least privileged)

MEMOS-2: Memory Segmentation in PM

Example

CS and DS
When in kernel
space

DS

CS

GDTR

GDT

User Code/Data

Drivers
Code/Data

Kernel Data

Kernel Code

Main Memory

GDT/LDT are stored in RAM

Up to 8192 segments

Segments can overlap

Can cover up to 4GB

Each GDT entry specifies

- Base address
- Size
- Growing direction
- Access rights:
 - Ring/Privilege
 - Executable
 - Read Access
 - Write Access

MEMOS-2: What does GRUB do for US?

- Switches to Protected Mode with a 4-GB flat memory segmentation
 - Our kernel can execute, read and write code and data anywhere in the first 4-GB
 - We are in Ring-0 and can run all the privileged instructions
- Finds our kernel in the disk and loads it at 0x100000 in the mem.
- Looks at the Multiboot header, gathers the information required to run the kernel according to the header
- Runs the kernel and passes the boot information according to the Multiboot Standard

MEMOS-2: The MultiBoot standard

- Kernel must define a header early on in its binary file in order to:
 - Specify what kind of information the bootloader must pass to the kernel
 - To verify if the binary file is a valid Multiboot-compliant kernel
- Defines the desired machine state before calling the kernel
- Defines the boot information format
 - It's a data structure that GRUB fills its fields that describe the system
 - Can you find the memory map there?
 - It's address in the memory is written in %EBX before jumping to the kernel
 - So the kernel can read EBX and access the information

MEMOS-2: C, sweet C!

- To call a C function:
 - Setup the stack pointer: %esp
 - Follow the [C calling convention](#) for argument passing
 - Jump to the function using the CALL instruction
- Make sure your C code does not depend on any external or GCC built-in library
 - Compiler flags: -fno-builtin -nostdinc

MEMOS-2: Video RAM

- No BIOS INT 0x10 to deal with the graphics ☹, but:
- You can ask GRUB (in the multiboot header) to set a specific video mode
- The basic 80x25 text-based VGA buffer is mapped to 0xB8000 in the main memory
- Changing each word of that buffer directly affects what gets displayed
- Maybe get a pointer and start poking at it?
- More info [here](#)

MEMOS-2: Example PM program (stub.S)

```
.text
.globl _start
_start:
    jmp     real_start
    # Multiboot header - Must be in 1st page of memory for GRUB
    .align 4
    .long   0x1BADB002 # Multiboot magic number
    .long   0x00000003 # Align modules to 4KB, req. mem size
    # See 'info multiboot' for further info
    .long   0xE4524FFB # Checksum

real_start:
    #TODO: Setup a proper stack for C
    #TODO: Prepare the boot information to pass to kmain
    call    kmain
    hlt
```


MEMOS-2: Example PM program (kentry.c)

```
#include "memos.h"
static unsigned short *videoram = (unsigned short *)0xB8000; //Base address of the VGA frame buffer
static int attrib = 0x0F; //black background, white foreground
static int csr_x = 0, csr_y = 0;
#define COLS 80

void putc(unsigned char c){
    if(c == 0x09){ // Tab (move to next multiple of 8)
        csr_x = (csr_x + 8) & ~(8 - 1);
    }else if(c == '\r'){ // CR
        csr_x = 0;
    }else if(c == '\n'){ // LF (unix-like)
        csr_x = 0; csr_y++;
    }else if(c >= ' '){ // Printable characters
        *(videoram + (csr_y * COLS + csr_x)) = c | (attrib << 8); // Put the character w/ attributes
        csr_x++;
    }
    if(csr_x >= COLS){ csr_x = 0; csr_y++;} // wrap around!
}

void puts(char *text){
    for (int i = 0; i < strlen((const char*)text); i++) // You know how to implement strlen ;)
        putc(text[i]);
}

void kmain(boot_info_t* binfo){
    puts("MemOS: Welcome *** Total Free Memory: ");
}
```


MEMOS-2 Example PM program (link.ld)

```
OUTPUT_FORMAT("elf32-i386")
OUTPUT_ARCH(i386)
ENTRY(_start)

SECTIONS {
    . = 0x100000;
    .ksection : {
        *(.*);
        . = ALIGN(0x1000);
    }
}
```


MEMOS-2: Example PM program (Makefile)

```
as --32 stub.S -o stub.o
```

```
gcc -m32 -fno-stack-protector -fno-builtin -nostdinc  
-c kentry.c -o kentry.o
```

```
ld -m elf_i386 -T link.ld stub.o kentry.o -o  
memos2.elf
```


MEMOS-2: Configuring GRUB

- Install grub on your disk image
- Copy your ELF binary image to the disk image
- Configure grub (through menu.lst) to load your ELF binary as a kernel

```
title MEMOS-2
```

```
root (hd0,0)
```

```
kernel /path/to/memos2.elf
```